

Nhom09.part1

July 11, 2021

BÀI TẬP LỚN MÔN HỌC MACHINE LEARNING K3- KHOA HỌC DỮ LIỆU

Nhóm 9:

Nguyễn Văn Thế

Lê Trung Hiếu

Nguyễn Thị Kim Duyên

1 Introduction

Bộ dữ liệu Stroke Prediction:

1. Thu thập dữ liệu:

- Theo Tổ chức Y tế Thế giới (WHO) đột quỵ là nguyên nhân gây tử vong thứ 2 trên toàn cầu, chiếm khoảng 11% tổng số ca tử vong. Bộ dữ liệu này được sử dụng để dự đoán liệu một bệnh nhân có khả năng bị đột quỵ hay không dựa trên các thông số đầu vào như giới tính, tuổi tác, các bệnh khác nhau và tình trạng hút thuốc. Mỗi hàng trong bộ dữ liệu cung cấp thông tin liên quan về bệnh nhân.
- Bộ dữ liệu được thu thập từ trang web: <https://www.kaggle.com/fedesoriano/stroke-prediction-dataset>
- Bộ dữ liệu gồm 5110 quan sát với 12 thuộc tính

2. Các trường dữ liệu:

- id:
 - Là số nhận dạng (ID) của bệnh nhân
 - Là biến định lượng liên tục kiểu số
- gender:
 - Giới tính bệnh nhân
 - Là biến định tính rời rạc kiểu String
 - Nhận 1 trong các giá trị: “Male”, “Female” or “Other”
- age:

- Tuổi của bệnh nhân
- Là biến định lượng liên tục kiểu số
- hypertension:
 - Chứng cao huyết áp của bệnh nhân
 - Là biến định tính rời rạc kiểu số
 - Nhận 1 trong 2 giá trị: 0 nếu bệnh nhân không bị cao huyết áp và 1 nếu bệnh nhân mắc cao huyết áp
- heart_disease:
 - Tình trạng đau tim của bệnh nhân
 - Là biến định tính rời rạc kiểu số
 - Nhận 1 trong 2 giá trị: 0 nếu bệnh nhân không bị đau tim, và 1 nếu bệnh nhân bị đau tim
- ever_married:
 - Tình trạng hôn nhân của bệnh nhân
 - Là biến định tính rời rạc kiểu String
 - Nhận 1 trong các giá trị: “No” nếu chưa kết hôn hoặc “Yes” nếu đã kết hôn
- work_type:
 - Loại nghề nghiệp của bệnh nhân
 - Là biến định tính rời rạc kiểu String
 - Nhận 1 trong các giá trị: “children”, “Govt_jov”, “Never_worked”, “Private” hoặc “Self-employed”
- Residence_type:
 - Nơi cư trú của bệnh nhân
 - Là biến định tính rời rạc kiểu String
 - Nhận 1 trong các giá trị: “Rural” hoặc “Urban”
- avg_glucose_level:
 - Lượng đường trung bình trong máu của bệnh nhân
 - Là biến định lượng liên tục kiểu số
- bmi:
 - Chỉ số đo lường cơ thể BMI của bệnh nhân
 - Là biến định lượng liên tục kiểu số
- smoking_status:
 - Tình trạng hút thuốc của bệnh nhân
 - Là biến định tính rời rạc kiểu String
 - Nhận 1 trong các giá trị: “formerly smoked”, “never smoked”, “smokes” or “Unknown”
- stroke:
 - Tình trạng đột quỵ của bệnh nhân
 - Là biến định tính rời rạc kiểu số

- Nhận 1 trong 2 giá trị: 0 nếu bệnh nhân không bị đột quy, và 1 nếu bệnh nhân bị đột quy

3. Mục đích phân tích dữ liệu:

Dự đoán đột quy ở con người Bị đột quỵ (mục tiêu phân tích, dự đoán) là biến “stroke” * Tình trạng đột quy của bệnh nhân * Là biến định tính rời rạc kiểu số * Nhận 1 trong 2 giá trị: 0 nếu bệnh nhân không bị đột quy, và 1 nếu bệnh nhân bị đột quy

1.1 Libraries and Utilities

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import itertools
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.ticker as mtick
from matplotlib.colors import ListedColormap, LinearSegmentedColormap

%matplotlib inline
import matplotlib

#Common model helpers
from sklearn.preprocessing import (StandardScaler,
                                   LabelEncoder,
                                   OneHotEncoder)

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.pipeline import Pipeline
from sklearn.tree import DecisionTreeRegressor
```

1.2 Data preprocessing

1.2.1 Loading data

```
[ ]: df = pd.read_csv('/content/drive/MyDrive/ML&DM/healthcare-dataset-stroke-data.
    ↪ csv')
```

```
[ ]: df.head(10).T
```

```
[ ]:
      id      0      1  ...      8      9
gender  9046  51676  ...  27419  60491
age      67      61  ...      59      78
```

| | | | | | |
|-------------------|-----------------|---------------|-----|---------|---------|
| hypertension | 0 | 0 | ... | 0 | 0 |
| heart_disease | 1 | 0 | ... | 0 | 0 |
| ever_married | Yes | Yes | ... | Yes | Yes |
| work_type | Private | Self-employed | ... | Private | Private |
| Residence_type | Urban | Rural | ... | Rural | Urban |
| avg_glucose_level | 228.69 | 202.21 | ... | 76.15 | 58.57 |
| bmi | 36.6 | NaN | ... | NaN | 24.2 |
| smoking_status | formerly smoked | never smoked | ... | Unknown | Unknown |
| stroke | 1 | 1 | ... | 1 | 1 |

[12 rows x 10 columns]

```
[ ]: # drop id column because it's unnecessary
df.drop('id',axis=1,inplace=True)
```

Attribute Information 1. **id**: unique identifier 2. **gender**: “Male”, “Female” or “Other” 3. **age**: age of the patient 4. **hypertension**: 0 if the patient doesn’t have hypertension, 1 if the patient has hypertension 5. **heart_disease**: 0 if the patient doesn’t have any heart diseases, 1 if the patient has a heart disease 6. **ever_married**: “No” or “Yes” 7. **work_type**: “children”, “Govt_jov”, “Never_worked”, “Private” or “Self-employed” 8. **Residence_type**: “Rural” or “Urban” 9. **avg_glucose_level**: average glucose level in blood 10. **bmi**: body mass index 11. **smoking_status**: “formerly smoked”, “never smoked”, “smokes” or “Unknown” 12. **stroke**: 1 if the patient had a stroke or 0 if not Note: “Unknown” in smoking_status means that the information is unavailable for this patient

```
[ ]: df.shape
```

```
[ ]: (5110, 11)
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                5110 non-null   object
1   age                   5110 non-null   float64
2   hypertension          5110 non-null   int64
3   heart_disease         5110 non-null   int64
4   ever_married          5110 non-null   object
5   work_type             5110 non-null   object
6   Residence_type        5110 non-null   object
7   avg_glucose_level     5110 non-null   float64
8   bmi                   4909 non-null   float64
9   smoking_status        5110 non-null   object
10  stroke                5110 non-null   int64
dtypes: float64(3), int64(3), object(5)
```

memory usage: 439.3+ KB

```
[ ]: df.describe()
```

```
[ ]:
count    age  hypertension  ...      bmi      stroke
mean     43.226614      0.097456  ...  28.893237    0.048728
std      22.612647      0.296607  ...   7.854067    0.215320
min       0.080000      0.000000  ...  10.300000    0.000000
25%      25.000000      0.000000  ...  23.500000    0.000000
50%      45.000000      0.000000  ...  28.100000    0.000000
75%      61.000000      0.000000  ...  33.100000    0.000000
max      82.000000      1.000000  ...  97.600000    1.000000
```

[8 rows x 6 columns]

```
[ ]: df.isnull().sum()
```

```
[ ]: gender          0
age                0
hypertension       0
heart_disease      0
ever_married       0
work_type          0
Residence_type     0
avg_glucose_level  0
bmi               201
smoking_status     0
stroke             0
dtype: int64
```

```
[ ]: # handling missing values
df['bmi'] = df['bmi'].fillna(df['bmi'].median())
df.isnull().sum()
```

```
[ ]: gender          0
age                0
hypertension       0
heart_disease      0
ever_married       0
work_type          0
Residence_type     0
avg_glucose_level  0
bmi                0
smoking_status     0
stroke             0
dtype: int64
```

1.2.2 Statistics of Categorical and Numerical Data

```
[ ]: # stats of numerical data
round (df.describe(exclude = 'object'), 2)
```

```
[ ]:
```

| | age | hypertension | ... | bmi | stroke |
|-------|---------|--------------|-----|---------|---------|
| count | 5110.00 | 5110.0 | ... | 5110.00 | 5110.00 |
| mean | 43.23 | 0.1 | ... | 28.86 | 0.05 |
| std | 22.61 | 0.3 | ... | 7.70 | 0.22 |
| min | 0.08 | 0.0 | ... | 10.30 | 0.00 |
| 25% | 25.00 | 0.0 | ... | 23.80 | 0.00 |
| 50% | 45.00 | 0.0 | ... | 28.10 | 0.00 |
| 75% | 61.00 | 0.0 | ... | 32.80 | 0.00 |
| max | 82.00 | 1.0 | ... | 97.60 | 1.00 |

[8 rows x 6 columns]

```
[ ]: # stats of categorical data
round (df.describe(exclude = ['float', 'int64']),2)
```

```
[ ]:
```

| | gender | ever_married | work_type | Residence_type | smoking_status |
|--------|--------|--------------|-----------|----------------|----------------|
| count | 5110 | 5110 | 5110 | 5110 | 5110 |
| unique | 3 | 2 | 5 | 2 | 4 |
| top | Female | Yes | Private | Urban | never smoked |
| freq | 2994 | 3353 | 2925 | 2596 | 1892 |

1.2.3 Initial Insights About Dataset

1. Data from bmi feature is missing.
2. Both categorical and numerical features are present. >* **Categorical Features:** gender, ever_married, work_type, Residence_type, smoking_status >* **Binary Numerical Features:** hypertension, heart_disease, stroke >* **Continous Numerical Features:** age, avg_glucose_level, bmi

```
[ ]: cont_features= ['age', 'avg_glucose_level', 'bmi']
cat_features=
→ ['gender', 'hypertension', 'heart_disease', 'ever_married', 'work_type', 'Residence_type', 'smoking_status']
target='stroke'
```

body mass index binning: <https://www.medicalnewstoday.com/articles/323446#body-mass-index>
Age binning: <https://kidspicturedictionary.com/english-through-pictures/people-english-through-pictures/age-physical-description/> average glucose binning: <https://agamatrix.com/blog/normal-blood-sugar-level-chart/>

```
[ ]: ## binning of numerical variables

df['bmi_cat'] = pd.cut(df['bmi'], bins = [0, 19, 25,30,10000], labels =
→ ['Underweight', 'Ideal', 'Overweight', 'Obesity'])
```

```
df['age_cat'] = pd.cut(df['age'], bins = [0,13,18, 45,60,200], labels =
↳['Children', 'Teens', 'Adults','Mid Adults','Elderly'])
df['glucose_cat'] = pd.cut(df['avg_glucose_level'], bins = [0,90,160,230,500],
↳labels = ['Low', 'Normal', 'High', 'Very High'])
```

2 Exploring Data Analysis

2.1 Distribution of Targets

```
[ ]: # Number of each category of the target label
df['stroke'].value_counts()
```


```
[ ]: 0    4861
     1     249
     Name: stroke, dtype: int64
```

```
[ ]: # Which category of the target label is how many percentage.
df['stroke'].value_counts()/len(df)*100
```

```
[ ]: 0    95.127202
     1     4.872798
     Name: stroke, dtype: float64
```

```
[ ]: sns.countplot(x = "stroke", data = df)
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8e92a35250>
```



2.1.1 Let's look at the numeric/continuous variable distribution

```
[ ]: variables = [variable for variable in df.columns if variable not in ['stroke']]

conts = ['age', 'avg_glucose_level', 'bmi']

[ ]: fig = plt.figure(figsize=(12, 12), dpi=150, facecolor='#fafafa')
gs = fig.add_gridspec(4, 3)
gs.update(wspace=0.1, hspace=0.4)

background_color = "#fafafa"


plot = 0
for row in range(0, 1):
    for col in range(0, 3):
        locals()["ax"+str(plot)] = fig.add_subplot(gs[row, col])
        locals()["ax"+str(plot)].set_facecolor(background_color)
        locals()["ax"+str(plot)].tick_params(axis='y', left=False)
        locals()["ax"+str(plot)].get_yaxis().set_visible(False)
        for s in ["top", "right", "left"]:
            locals()["ax"+str(plot)].spines[s].set_visible(False)
        plot += 1

plot = 0
for variable in conts:
    sns.kdeplot(df[variable], ax=locals()["ax"+str(plot)], color='#0f4c81',
    →shade=True, linewidth=1.5, ec='black', alpha=0.9, zorder=3, legend=False)
    locals()["ax"+str(plot)].grid(which='major', axis='x', zorder=0,
    →color='gray', linestyle=':', dashes=(1,5))
    #locals()["ax"+str(plot)].set_xlabel(variable) removed this for
    →aesthetics
    plot += 1

ax0.set_xlabel('Age')
ax1.set_xlabel('Avg. Glucose Levels')
ax2.set_xlabel('BMI')

ax0.text(-20, 0.022, 'Numeric Variable Distribution', fontsize=13,
    →fontweight='bold', fontfamily='serif')
ax0.text(-20, 0.02, 'We see a positive skew in BMI and Glucose Level',
    →fontsize=11, fontweight='light', fontfamily='serif')

plt.show()
```

Let's see how the distribution of our numeric variables is different for those that have strokes, and those that do not.

```
[ ]: fig = plt.figure(figsize=(12, 12), dpi=150, facecolor=background_color)
gs = fig.add_gridspec(4, 3)
gs.update(wspace=0.1, hspace=0.4)

plot = 0
for row in range(0, 1):
    for col in range(0, 3):
        locals()["ax"+str(plot)] = fig.add_subplot(gs[row, col])
        locals()["ax"+str(plot)].set_facecolor(background_color)
        locals()["ax"+str(plot)].tick_params(axis='y', left=False)
        locals()["ax"+str(plot)].get_yaxis().set_visible(False)
        for s in ["top", "right", "left"]:
            locals()["ax"+str(plot)].spines[s].set_visible(False)
        plot += 1

plot = 0

s = df[df['stroke'] == 1]
ns = df[df['stroke'] == 0]

for feature in conts:
    sns.kdeplot(s[feature], ax=locals()["ax"+str(plot)], color='#0f4c81',
    ↪shade=True, linewidth=1.5, ec='black', alpha=0.9, zorder=3, legend=False)
    sns.kdeplot(ns[feature], ax=locals()["ax"+str(plot)], color='#9bb7d4',
    ↪shade=True, linewidth=1.5, ec='black', alpha=0.9, zorder=3, legend=False)
    locals()["ax"+str(plot)].grid(which='major', axis='x', zorder=0,
    ↪color='gray', linestyle=':', dashes=(1,5))
    #locals()["ax"+str(plot)].set_xlabel(feature)
    plot += 1
```

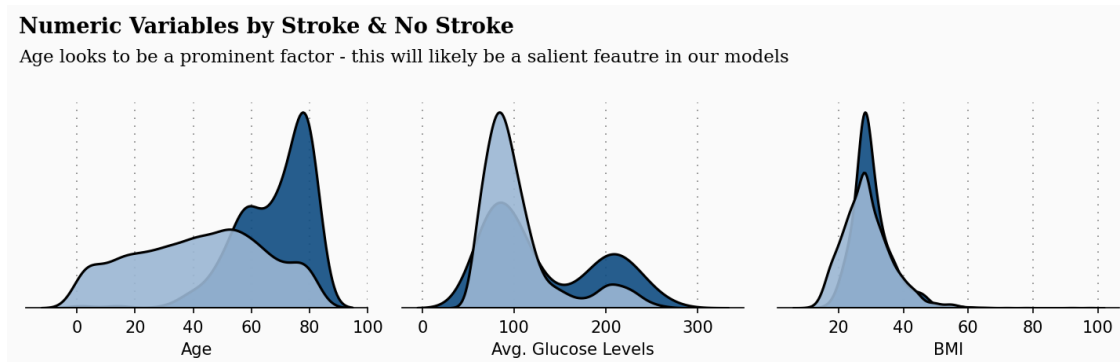
```

ax0.set_xlabel('Age')
ax1.set_xlabel('Avg. Glucose Levels')
ax2.set_xlabel('BMI')

ax0.text(-20, 0.056, 'Numeric Variables by Stroke & No Stroke', fontsize=13,
        fontweight='bold', fontfamily='serif')
ax0.text(-20, 0.05, 'Age looks to be a prominent factor - this will likely be a
        salient feautre in our models',
        fontsize=11, fontweight='light', fontfamily='serif')

plt.show()

```



Based on the above plots, it seems clear that Age is a big factor in stroke patients - the older you get the more at risk

```

[ ]: str_only = df[df['stroke'] == 1]
     no_str_only = df[df['stroke'] == 0]

```

```

[ ]: # Setting up figure and axes

```

```

fig = plt.figure(figsize=(10,16),dpi=150,facecolor=background_color)
gs = fig.add_gridspec(4, 2)
gs.update(wspace=0.5, hspace=0.2)
ax0 = fig.add_subplot(gs[0, 0:2])
ax1 = fig.add_subplot(gs[1, 0:2])

ax0.set_facecolor(background_color)
ax1.set_facecolor(background_color)

# glucose

sns.regplot(no_str_only['age'],y=no_str_only['avg_glucose_level'],
            color='lightgray',
            logx=True,

```

```

        ax=ax0)

sns.regplot(str_only['age'],y=str_only['avg_glucose_level'],
            color='#0f4c81',
            logx=True,scatter_kws={'edgecolors':['black'],
                                   'linewidth': 1},

            ax=ax0)

ax0.set(ylim=(0, None))
ax0.set_xlabel(" ",fontsize=12,fontfamily='serif')
ax0.set_ylabel("Avg. Glucose Level",fontsize=10,fontfamily='serif')

ax0.tick_params(axis='x', bottom=False)
ax0.get_xaxis().set_visible(False)

for s in ['top','left','bottom']:
    ax0.spines[s].set_visible(False)

# bmi
sns.regplot(no_str_only['age'],y=no_str_only['bmi'],
            color='lightgray',
            logx=True,
            ax=ax1)

sns.regplot(str_only['age'],y=str_only['bmi'],
            color='#0f4c81', scatter_kws={'edgecolors':['black'],
                                   'linewidth': 1},

            logx=True,
            ax=ax1)

ax1.set_xlabel("Age",fontsize=10,fontfamily='serif')
ax1.set_ylabel("BMI",fontsize=10,fontfamily='serif')

for s in ['top','left','right']:
    ax0.spines[s].set_visible(False)
    ax1.spines[s].set_visible(False)

ax0.text(-5,350,'Strokes by Age, Glucose Level, and
↳BMI',fontsize=18,fontfamily='serif',fontweight='bold')
ax0.text(-5,320,'Age appears to be a very important
↳factor',fontsize=14,fontfamily='serif')

ax0.tick_params(axis=u'both', which=u'both',length=0)

```

```
ax1.tick_params(axis=u'both', which=u'both',length=0)
```

```
plt.show()
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.


FutureWarning

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Strokes by Age, Glucose Level, and BMI

Age appears to be a very important factor



Age is a big factor, and also has slight relationships with BMI & Avg. Glucose levels.

```
[ ]: fig = plt.figure(figsize=(10, 5), dpi=150, facecolor=background_color)
gs = fig.add_gridspec(2, 1)
gs.update(wspace=0.11, hspace=0.5)
ax0 = fig.add_subplot(gs[0, 0])
ax0.set_facecolor(background_color)

df['age'] = df['age'].astype(int)

rate = []
for i in range(df['age'].min(), df['age'].max()):
    rate.append(df[df['age'] < i]['stroke'].sum() / len(df[df['age'] <
    ↪ i]['stroke']))

sns.lineplot(data=rate, color='#0f4c81', ax=ax0)
```

```

for s in ["top","right","left"]:
    ax0.spines[s].set_visible(False)

ax0.tick_params(axis='both', which='major', labelsize=8)
ax0.tick_params(axis=u'both', which=u'both',length=0)

ax0.text(-3,0.055,'Risk Increase by_
→Age',fontsize=18,fontfamily='serif',fontweight='bold')
ax0.text(-3,0.047,'As age increase, so too does risk of having a_
→stroke',fontsize=14,fontfamily='serif')


plt.show()

```

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:12: RuntimeWarning:
invalid value encountered in long_scalars
  if sys.path[0] == '':

```



```

[ ]: # Drop single 'Other' gender
no_str_only = no_str_only[(no_str_only['gender'] != 'Other')]

```

```

[ ]: fig = plt.figure(figsize=(22,15))
gs = fig.add_gridspec(3, 3)
gs.update(wspace=0.35, hspace=0.27)
ax0 = fig.add_subplot(gs[0, 0])
ax1 = fig.add_subplot(gs[0, 1])
ax2 = fig.add_subplot(gs[0, 2])
ax3 = fig.add_subplot(gs[1, 0])
ax4 = fig.add_subplot(gs[1, 1])
ax5 = fig.add_subplot(gs[1, 2])
ax6 = fig.add_subplot(gs[2, 0])
ax7 = fig.add_subplot(gs[2, 1])
ax8 = fig.add_subplot(gs[2, 2])

background_color = "#f6f6f6"

```

```

fig.patch.set_facecolor(background_color) # figure background color

# Plots

## Age

ax0.grid(color='gray', linestyle=':', axis='y', zorder=0, dashes=(1,5))
positive = pd.DataFrame(str_only["age"])
negative = pd.DataFrame(no_str_only["age"])
sns.kdeplot(positive["age"], ax=ax0,color="#0f4c81", shade=True,
    ↳ec='black',label="positive")
sns.kdeplot(negative["age"], ax=ax0, color="#9bb7d4", shade=True,
    ↳ec='black',label="negative")
#ax3.text(0.29, 13, 'Age',
#    ↳fontsize=14, fontweight='bold', fontfamily='serif', color="#323232")
ax0.yaxis.set_major_locator(mtick.MultipleLocator(2))
ax0.set_ylabel('')
ax0.set_xlabel('')
ax0.text(-20, 0.0465, 'Age', fontsize=14, fontweight='bold',
    ↳fontfamily='serif', color="#323232")

# Smoking
positive = pd.DataFrame(str_only["smoking_status"].value_counts())
positive["Percentage"] = positive["smoking_status"].apply(lambda x: x/
    ↳sum(positive["smoking_status"])*100)
negative = pd.DataFrame(no_str_only["smoking_status"].value_counts())
negative["Percentage"] = negative["smoking_status"].apply(lambda x: x/
    ↳sum(negative["smoking_status"])*100)

ax1.text(0, 4, 'Smoking Status', fontsize=14, fontweight='bold',
    ↳fontfamily='serif', color="#323232")
ax1.barh(positive.index, positive['Percentage'], color="#0f4c81", zorder=3,
    ↳height=0.7)
ax1.barh(negative.index, negative['Percentage'], color="#9bb7d4",
    ↳zorder=3,ec='black', height=0.3)
ax1.xaxis.set_major_formatter(mtick.PercentFormatter())
ax1.xaxis.set_major_locator(mtick.MultipleLocator(10))

##
# Ax2 - GENDER
positive = pd.DataFrame(str_only["gender"].value_counts())

```

```

positive["Percentage"] = positive["gender"].apply(lambda x: x/
↳sum(positive["gender"])*100)
negative = pd.DataFrame(no_str_only["gender"].value_counts())
negative["Percentage"] = negative["gender"].apply(lambda x: x/
↳sum(negative["gender"])*100)

x = np.arange(len(positive))
ax2.text(-0.4, 68.5, 'Gender', fontsize=14, fontweight='bold',
↳fontfamily='serif', color="#323232")
ax2.grid(color='gray', linestyle=':', axis='y', zorder=0, dashes=(1,5))
ax2.bar(x, height=positive["Percentage"], zorder=3, color="#0f4c81", width=0.4)
ax2.bar(x+0.4, height=negative["Percentage"], zorder=3, color="#9bb7d4",
↳width=0.4)
ax2.set_xticks(x + 0.4 / 2)
ax2.set_xticklabels(['Male', 'Female'])
ax2.yaxis.set_major_formatter(mtick.PercentFormatter())
ax2.yaxis.set_major_locator(mtick.MultipleLocator(10))
for i,j in zip([0, 1], positive["Percentage"]):
    ax2.annotate(f'{j:0.0f}%', xy=(i, j/2), color='#f6f6f6',
↳horizontalalignment='center', verticalalignment='center')
for i,j in zip([0, 1], negative["Percentage"]):
    ax2.annotate(f'{j:0.0f}%', xy=(i+0.4, j/2), color='#f6f6f6',
↳horizontalalignment='center', verticalalignment='center')

##
# Ax9 - residence type
positive = pd.DataFrame(str_only["rece"].value_counts())
positive["Percentage"] = positive["gender"].apply(lambda x: x/
↳sum(positive["gender"])*100)
negative = pd.DataFrame(no_str_only["gender"].value_counts())
negative["Percentage"] = negative["gender"].apply(lambda x: x/
↳sum(negative["gender"])*100)

x = np.arange(len(positive))
ax2.text(-0.4, 68.5, 'Gender', fontsize=14, fontweight='bold',
↳fontfamily='serif', color="#323232")
ax2.grid(color='gray', linestyle=':', axis='y', zorder=0, dashes=(1,5))
ax2.bar(x, height=positive["Percentage"], zorder=3, color="#0f4c81", width=0.4)
ax2.bar(x+0.4, height=negative["Percentage"], zorder=3, color="#9bb7d4",
↳width=0.4)
ax2.set_xticks(x + 0.4 / 2)
ax2.set_xticklabels(['Male', 'Female'])
ax2.yaxis.set_major_formatter(mtick.PercentFormatter())
ax2.yaxis.set_major_locator(mtick.MultipleLocator(10))
for i,j in zip([0, 1], positive["Percentage"]):

```



```

        ax2.annotate(f'{j:0.0f}%',xy=(i, j/2), color='#f6f6f6',
        ↪horizontalalignment='center', verticalalignment='center')
for i,j in zip([0, 1], negative["Percentage"]):
    ax2.annotate(f'{j:0.0f}%',xy=(i+0.4, j/2), color='#f6f6f6',
    ↪horizontalalignment='center', verticalalignment='center')

# Heart Dis

positive = pd.DataFrame(str_only["heart_disease"].value_counts())
positive["Percentage"] = positive["heart_disease"].apply(lambda x: x/
    ↪sum(positive["heart_disease"])*100)
negative = pd.DataFrame(no_str_only["heart_disease"].value_counts())
negative["Percentage"] = negative["heart_disease"].apply(lambda x: x/
    ↪sum(negative["heart_disease"])*100)

x = np.arange(len(positive))
ax3.text(-0.3, 110, 'Heart Disease', fontsize=14, fontweight='bold',
    ↪fontfamily='serif', color="#323232")
ax3.grid(color='gray', linestyle=':', axis='y', zorder=0, dashes=(1,5))
ax3.bar(x, height=positive["Percentage"], zorder=3, color="#0f4c81", width=0.4)
ax3.bar(x+0.4, height=negative["Percentage"], zorder=3, color="#9bb7d4",
    ↪width=0.4)
ax3.set_xticks(x + 0.4 / 2)
ax3.set_xticklabels(['No History', 'History'])
ax3.yaxis.set_major_formatter(mtick.PercentFormatter())
ax3.yaxis.set_major_locator(mtick.MultipleLocator(20))
for i,j in zip([0, 1], positive["Percentage"]):
    ax3.annotate(f'{j:0.0f}%',xy=(i, j/2), color='#f6f6f6',
    ↪horizontalalignment='center', verticalalignment='center')
for i,j in zip([0, 1], negative["Percentage"]):
    ax3.annotate(f'{j:0.0f}%',xy=(i+0.4, j/2), color='#f6f6f6',
    ↪horizontalalignment='center', verticalalignment='center')

## AX4 - TITLE

ax4.spines["bottom"].set_visible(False)
ax4.tick_params(left=False, bottom=False)
ax4.set_xticklabels([])
ax4.set_yticklabels([])
# ax4.text(0.5, 0.6, 'Can we see patterns for\n\n patients in our data?',
    ↪horizontalalignment='center', verticalalignment='center',
#         fontsize=22, fontweight='bold', fontfamily='serif', color="#323232")

```

```

ax4.text(0.15,0.57,"Stroke", fontweight="bold", fontfamily='serif',
    ↳ fontsize=22, color='#0f4c81')
ax4.text(0.41,0.57,"&", fontweight="bold", fontfamily='serif', fontsize=22,
    ↳ color='#323232')
ax4.text(0.49,0.57,"No-Stroke", fontweight="bold", fontfamily='serif',
    ↳ fontsize=22, color='#9bb7d4')

# Glucose

ax5.grid(color='gray', linestyle=':', axis='y', zorder=0, dashes=(1,5))
positive = pd.DataFrame(str_only["avg_glucose_level"])
negative = pd.DataFrame(no_str_only["avg_glucose_level"])
sns.kdeplot(positive["avg_glucose_level"], ax=ax5,color="#0f4c81",ec='black',
    ↳ shade=True, label="positive")
sns.kdeplot(negative["avg_glucose_level"], ax=ax5, color="#9bb7d4",
    ↳ ec='black',shade=True, label="negative")
ax5.text(-55, 0.01855, 'Avg. Glucose Level',
    ↳ fontsize=14, fontweight='bold', fontfamily='serif', color="#323232")
ax5.yaxis.set_major_locator(mtick.MultipleLocator(2))
ax5.set_ylabel('')
ax5.set_xlabel('')

## BMI

ax6.grid(color='gray', linestyle=':', axis='y', zorder=0, dashes=(1,5))
positive = pd.DataFrame(str_only["bmi"])
negative = pd.DataFrame(no_str_only["bmi"])
sns.kdeplot(positive["bmi"], ax=ax6,color="#0f4c81", ec='black',shade=True,
    ↳ label="positive")
sns.kdeplot(negative["bmi"], ax=ax6, color="#9bb7d4",ec='black', shade=True,
    ↳ label="negative")
ax6.text(-0.06, 0.09, 'BMI',
    ↳ fontsize=14, fontweight='bold', fontfamily='serif', color="#323232")
ax6.yaxis.set_major_locator(mtick.MultipleLocator(2))
ax6.set_ylabel('')
ax6.set_xlabel('')

# Work Type

positive = pd.DataFrame(str_only["work_type"].value_counts())

```

```

positive["Percentage"] = positive["work_type"].apply(lambda x: x/
↳sum(positive["work_type"])*100)
positive = positive.sort_index()

negative = pd.DataFrame(no_str_only["work_type"].value_counts())
negative["Percentage"] = negative["work_type"].apply(lambda x: x/
↳sum(negative["work_type"])*100)
negative = negative.sort_index()

ax7.bar(negative.index, height=negative["Percentage"], zorder=3,
↳color="#9bb7d4", width=0.05)
ax7.scatter(negative.index, negative["Percentage"], zorder=3,s=200,
↳color="#9bb7d4")
ax7.bar(np.arange(len(positive.index))+0.4, height=positive["Percentage"],
↳zorder=3, color="#0f4c81", width=0.05)
ax7.scatter(np.arange(len(positive.index))+0.4, positive["Percentage"],
↳zorder=3,s=200, color="#0f4c81")

ax7.yaxis.set_major_formatter(mtick.PercentFormatter())
ax7.yaxis.set_major_locator(mtick.MultipleLocator(10))
ax7.set_xticks(np.arange(len(positive.index))+0.4 / 2)
ax7.set_xticklabels(list(positive.index),rotation=0)
ax7.text(-0.5, 66, 'Work Type', fontsize=14, fontweight='bold',
↳fontfamily='serif', color="#323232")

# hypertension

positive = pd.DataFrame(str_only["hypertension"].value_counts())
positive["Percentage"] = positive["hypertension"].apply(lambda x: x/
↳sum(positive["hypertension"])*100)
negative = pd.DataFrame(no_str_only["hypertension"].value_counts())
negative["Percentage"] = negative["hypertension"].apply(lambda x: x/
↳sum(negative["hypertension"])*100)

x = np.arange(len(positive))
ax8.text(-0.45, 100, 'Hypertension', fontsize=14, fontweight='bold',
↳fontfamily='serif', color="#323232")
ax8.grid(color='gray', linestyle=':', axis='y', zorder=0, dashes=(1,5))
ax8.bar(x, height=positive["Percentage"], zorder=3, color="#0f4c81", width=0.4)
ax8.bar(x+0.4, height=negative["Percentage"], zorder=3, color="#9bb7d4",
↳width=0.4)
ax8.set_xticks(x + 0.4 / 2)
ax8.set_xticklabels(['No History','History'])
ax8.yaxis.set_major_formatter(mtick.PercentFormatter())

```

```

ax8.yaxis.set_major_locator(mtick.MultipleLocator(20))
for i,j in zip([0, 1], positive["Percentage"]):
    ax8.annotate(f'{j:0.0f}%',xy=(i, j/2), color='#f6f6f6',
    ↪horizontalalignment='center', verticalalignment='center')
for i,j in zip([0, 1], negative["Percentage"]):
    ax8.annotate(f'{j:0.0f}%',xy=(i+0.4, j/2), color='#f6f6f6',
    ↪horizontalalignment='center', verticalalignment='center')

# tidy up

for s in ["top","right","left"]:
    for i in range(0,9):
        locals()["ax"+str(i)].spines[s].set_visible(False)

for i in range(0,9):
    locals()["ax"+str(i)].set_facecolor(background_color)
    locals()["ax"+str(i)].tick_params(axis=u'both', which=u'both',length=0)
    locals()["ax"+str(i)].set_facecolor(background_color)

plt.show()

```



```

[ ]: df_copy = df.copy()
      # feature log transformations

df_copy['age'] = df_copy['age'].apply(lambda x: np.log(x+10)*3)
df_copy['avg_glucose_level'] = df_copy['avg_glucose_level'].apply(lambda x: np.
    ↪ log(x+10)*2)
df_copy['bmi'] = df_copy['bmi'].apply(lambda x: np.log(x+10)*2)


# preprocessing - label encoding and numerical value scaling
ohe = OneHotEncoder()
ss = StandardScaler()
le = LabelEncoder()

## label encoding of ordinal categorical features
for col in df_copy.columns:
    df_copy[col] = le.fit_transform(df_copy[col])

cols = df_copy.columns
## normalizing with standard scaler of numerical features
df_copy[cols] = ss.fit_transform(df_copy[cols])


# correlation map for all the features
df_corr = df_copy.corr()
mask = np.triu(np.ones_like(df_corr, dtype=np.bool))

fig, ax = plt.subplots(figsize = (8,8))
fig.patch.set_facecolor('#f6f5f5')
ax.set_facecolor('#f6f5f5')

mask = mask[1:, :-1]
corr = df_corr.iloc[1:, :-1].copy()

colors = ['#fafafa', '#512b58', '#fe346e']
colormap = matplotlib.colors.LinearSegmentedColormap.from_list("", colors)

# plot heatmap
sns.heatmap(corr, mask=mask, annot=True, fmt=".2f", cmap = colormap,
    vmin=-0.15, vmax=0.5, cbar_kws={"shrink": .5, }, ax = ax, cbar =
    ↪ False,
    linewidth = 1, linecolor = '#f6f5f5', square = True, annot_kws =
    ↪ {'size':10, 'color':'black'} )
# yticks
ax.tick_params(axis = 'y', rotation=0)

```

```

xticks = ['Gender', 'Age', 'Hyper tension', 'Heart Disease', 'Marriage', 'Work',
↳ 'Residence', 'Glucose Level', 'BMI', 'Smoking Status', 'Stroke', 'BMI
↳ Cat', 'Age Cat']
yticks = ['Gender', 'Age', 'Hyper tension', 'Heart Disease', 'Marriage', 'Work',
↳ 'Residence', 'Glucose Level', 'BMI', 'Smoking Status', 'Stroke', 'BMI
↳ Cat', 'Age Cat']
ax.set_xticklabels(xticks, {'size':10, 'weight':'bold'}, rotation = 90, alpha =
↳ 0.9)
ax.set_yticklabels(yticks, {'size':10, 'weight':'bold'}, rotation = 0, alpha =
↳ 0.9)
ax.text(-3.5,-1.1, 'Correlation Map of Features',{'size': 16, 'weight':'bold'},
↳ alpha = 0.9)
ax.text(-3.5,-0.65, 'A Glipse on feature correlation for processed feature data.
↳ ',{'size': 12, 'weight':'normal'}, alpha = 0.8)

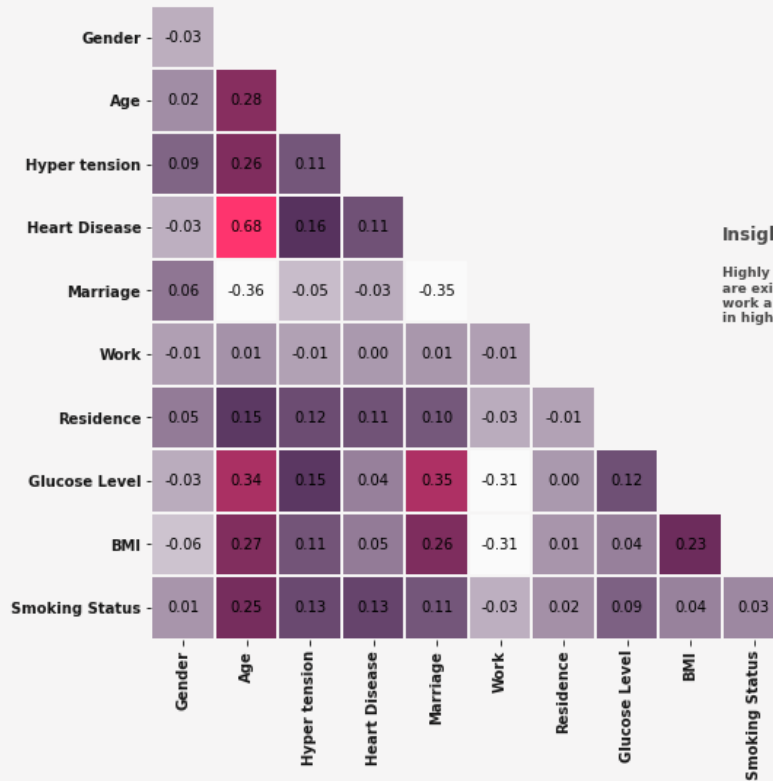
ax.text(9,5, 'Highly correlated positive correlations \nare exist for Age and
↳ Marriage, while \nwork and other are correlated negatively \nin highest
↳ order.',{'size': 9, 'weight':'bold'},alpha = 0.7)
ax.text(9,3.7, 'Insight:',{'size': 12, 'weight':'bold'},alpha = 0.7)

fig.show()

```

Correlation Map of Features

A Glipse on feature correlation for processed feature data.



Insight:

Highly correlated positive correlations are exist for Age and Marriage, while work and other are correlated negatively in highest order.

```
[ ]: labels = ['Smoking', 'BMI', 'Age', 'Marriage', 'Heart Disease',
    ↳ 'Stroke', 'Hypertension', 'Age Cat', 'Gender', 'Work', 'BMI Cat',
    ↳ 'Residence', 'Glucose Level', 'Glucose Cat']

g = sns.clustermap(df_corr, annot = True, fmt = '0.2f',
    cbar= False, cbar_pos=(0,0, 0,0),linewidth = 0.5,
    cmap = colormap,dendrogram_ratio=0.1,
    facecolor = '#f6f5f5', figsize = (8,8),square = True,
    annot_kws = {'size':10, 'color':'black'})

plt.gcf().set_facecolor('#f6f5f5')
label_args = {'font':18, 'weight':'bold'}
plt.setp(g.ax_heatmap.set_yticklabels(labels), rotation=0, fontsize = 10,
    ↳fontfamily = 'Serif', fontweight = 'bold', alpha = 0.8) # For y axis
plt.setp(g.ax_heatmap.set_xticklabels(labels), rotation=90, fontsize = 10,
    ↳fontfamily = 'Serif', fontweight = 'bold', alpha = 0.8) # For x axis
g.fig.text(0,1.065,'Visualization of Clustering of Each Feature with
    ↳Other',{'size':16, 'weight':'bold'})
```

```
g.fig.text(0,1.015,'Lines on the top and left of the cluster map are called_\ndendrograms, which indiate the dependency of features.',{'size':12}, alpha_\n→= 0.8)\nplt.show()
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/matrix.py:1216: UserWarning:\n``square=True`` ignored in clustermap\n  warnings.warn(msg)
```

